



ILLINOIS

UNIVERSITY OF ILLINOIS AT URBANA-CHAMPAIGN

CS498

Applied Machine Learning

Assignment #8

(Extra Homework for Four Hour version)

Students:

nidiaib2, Nidia Bucarelli

sunnyk2, Sunny Katiyar

wangx2, Wang Xiang

May 14, 2020

Spring 2020


Images were perturbed with the following script:

```
def fgsm_attack(image, epsilon, data_grad):  
    sign_data_grad = data_grad.sign()  
    perturbed_image = image + epsilon*sign_data_grad  
    perturbed_image = torch.clamp(perturbed_image, 0, 1)  
    return perturbed_image
```


Accuracy of classifying the next digit: **99.3%**

Example of outputs:

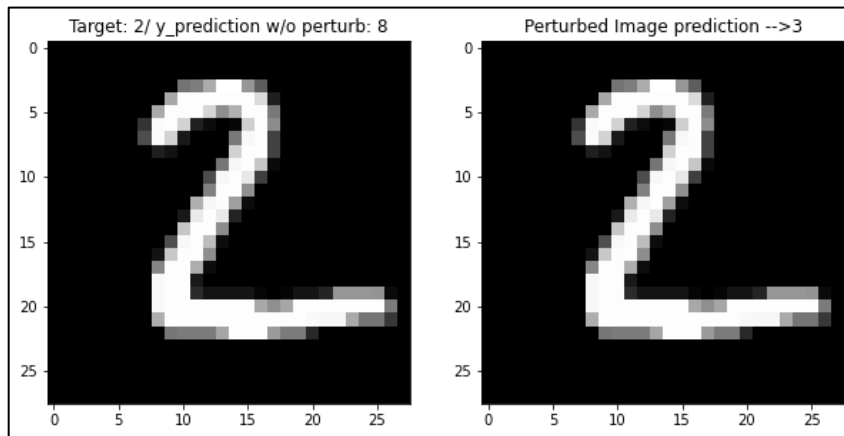
Digit-0




Digit-1



Digit-2




Digit-3



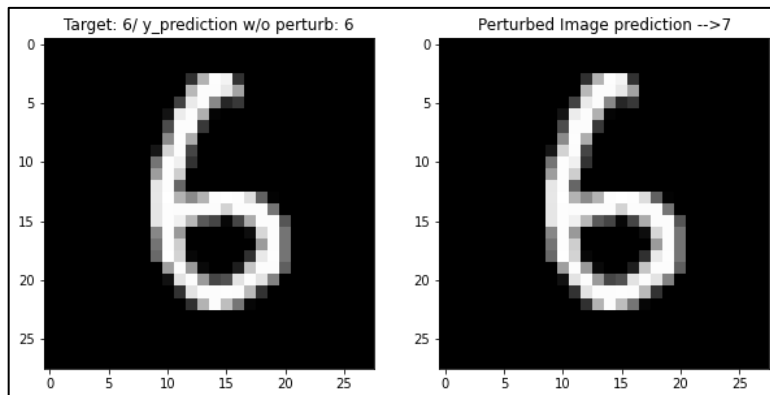
Digit-4



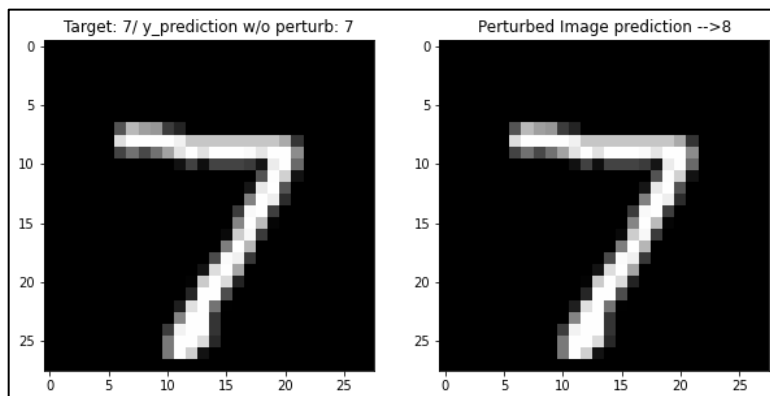
Digit-5



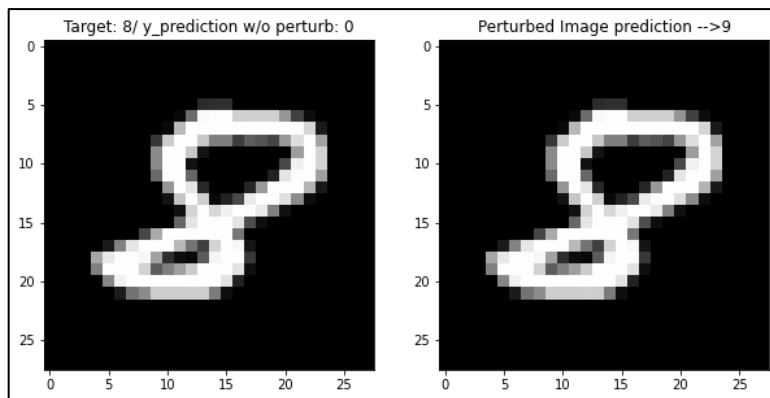
Digit-6




Digit-7



Digit-8



Digit-9



CS-498_Extra Homework

May 15, 2020

```
In [0]: import h5py
import numpy as np
from random import randint
import time
import requests
import matplotlib.pyplot as plt
import cv2
# from sklearn.ensemble.forest import _generate_unsampled_indices
# from sklearn.ensemble import BaggingClassifier
from sklearn.ensemble import RandomForestClassifier
import h5py
import numpy as np
from random import randint
import torch
import torch.nn as nn
import torchvision
import torchvision.transforms as transforms
import torch.nn.functional as F
import torch.optim as optim
from torch.autograd import Variable
import time
import os

##Load MNIST DATA

batch_size= 1
train_dataset= torchvision.datasets.MNIST(root= './', train= True, transform= transform)
test_dataset=torchvision.datasets.MNIST(root= './', transform= transforms.ToTensor(),

train_loader= torch.utils.data.DataLoader(dataset= train_dataset, batch_size=batch_size,
test_loader= torch.utils.data.DataLoader(dataset= test_dataset, batch_size=batch_size,

In [0]: os.chdir('/content/drive/My Drive/CS-498 Applied Machine Learning/HW8/Nidia')
# pretrained_model= torch.load('/content/drive/My Drive/CS-498 Applied Machine Learning

In [0]: targets= list(range(0,10))
target_new= []
```

```

l= []
for t in range (0, len(targets)):
    for k, y in train_loader:
        if y == targets[t]:
            target_new.append(k)
            l.append(y)
            break

```

In [0]: *##MODEL ARCHITECTURE*

```

class network(nn.Module):
    def __init__(self):
        super(network, self).__init__()
        self.conv1 = nn.Conv2d(1,20,5,stride=1,padding=0)
        # self.dropout1 = nn.Dropout(0.30)
        self.MaxPool1= nn.MaxPool2d(2, stride=2)

        self.conv2 = nn.Conv2d(20,50,5,stride=1, padding=0)
        # self.dropout2 = nn.Dropout(0.30)
        self.MaxPool2= nn.MaxPool2d(2, stride=2)

        self.conv3 = nn.Conv2d(50,500,4,stride=1, padding=0)
        # self.conv3_bn= nn.BatchNorm2d(num_features=500,track_running_stats=False)

        self.conv4 = nn.Conv2d(500,10,1,stride=1, padding=0)
        # self.conv4_bn= nn.BatchNorm2d(num_features=10,track_running_stats=False)

        # self.dropoutfc= nn.Dropout(0.30)
        # self.fc3 = nn.Linear(500,10)

    def forward(self, x):
        # x= self.MaxPool1(self.dropout1(self.conv1(x)))
        # x= self.MaxPool2(self.dropout2(self.conv2(x)))

        x= self.MaxPool1((self.conv1(x)))
        x= self.MaxPool2((self.conv2(x)))
        # x= self.conv3_bn((self.conv3(x)))
        x= self.conv3(x)
        x= F.relu(x)
        x= self.conv4(x)
        # x= self.conv4_bn(x)

        return (x.flatten()).reshape(x.shape[0],10))

```

In [0]: `def fgsm_attack(image, epsilon, data_grad):`
`sign_data_grad = data_grad.sign()`

```

    # Create the perturbed image by adjusting each pixel of the input image
    perturbed_image = image + epsilon*sign_data_grad
    # Adding clipping to maintain [0,1] range
    perturbed_image = torch.clamp(perturbed_image, 0, 1)
    # Return the perturbed image
    return perturbed_image

```

```

In [0]: device= torch.device('cuda' if torch.cuda.is_available() else 'cpu')
        model = network().to(device)

```

```

        checkpoint = torch.load('/content/drive/My Drive/CS-498 Applied Machine Learning/HW8/N
        model.load_state_dict(checkpoint['state_dict'])
        model.eval()

```

```

Out[0]: network(
  (conv1): Conv2d(1, 20, kernel_size=(5, 5), stride=(1, 1))
  (MaxPool1): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (conv2): Conv2d(20, 50, kernel_size=(5, 5), stride=(1, 1))
  (MaxPool2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (conv3): Conv2d(50, 500, kernel_size=(4, 4), stride=(1, 1))
  (conv4): Conv2d(500, 10, kernel_size=(1, 1), stride=(1, 1))
)

```

```

In [0]: criterion= nn.CrossEntropyLoss()
        learning_rate = 0.01
        optimizer= optim.SGD(model.parameters(), lr=learning_rate)
        adv_ex = []
        adv_examples= []
        num_epoch = 1
        # epsilons = [0, .05, .1, .15, .2, .25, .3]
        epsilons = 0.05
        num_epoch = 1
        for epoch in range (0, num_epoch):
            correct = 0
            adv_examples = []
            for data, target in test_loader:
                data, target = data.to(device), target.to(device)
                y_alternate = target + 1
                if y_alternate == 10:
                    y_alternate = torch.tensor([0]).to(device)
                y_alternate= y_alternate.to(device)

                # Set requires_grad attribute of tensor. Important for Attack
                data.requires_grad = True

                # Forward pass the data through the model
                output = model(data)
                y_prediction= output.data.max(1)[1]

```



```

# # If the initial prediction is wrong, dont bother attacking, just move on
if y_prediction == y_alternate:
    correct += 1
    continue

# Calculate the loss
loss = criterion(output, y_alternate)

model.zero_grad()
# Calculate gradients of model in backward pass
loss.backward()
optimizer.step()

# Collect datagrad
data_grad = data.grad.data

perturbed_data = fgsm_attack(data, epsilon, data_grad)

output = model(perturbed_data)
# Check for success
final_pred = output.max(1, keepdim=True)[1] # get the index of the max log-probability

if final_pred == y_alternate:
    correct = correct + 1
    if len(adv_examples) < 100:
        adv_ex = perturbed_data.squeeze().detach().cpu().numpy()
        original = data.squeeze().detach().cpu().numpy()
        adv_examples.append((target, y_prediction.item(), final_pred.item(), original))

final_acc = correct/float(len(test_loader))
print('epoch:', epoch, "Epsilon: {} \t Test Accuracy = {} / {} = {}".format(epsilon, c

```

```
epoch: 0 Epsilon: 0.05          Test Accuracy = 9931 / 10000 = 0.9931
```



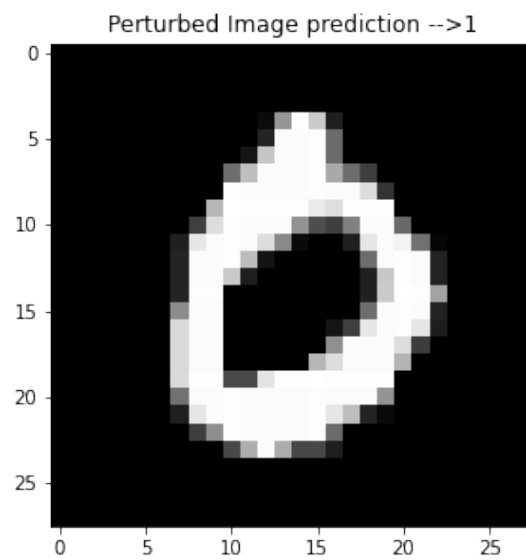

```


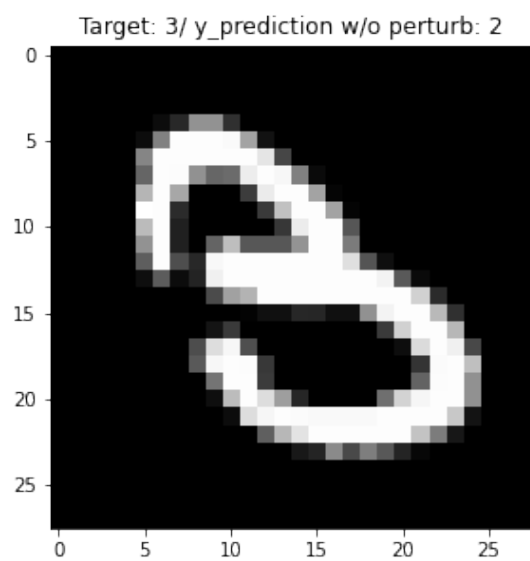
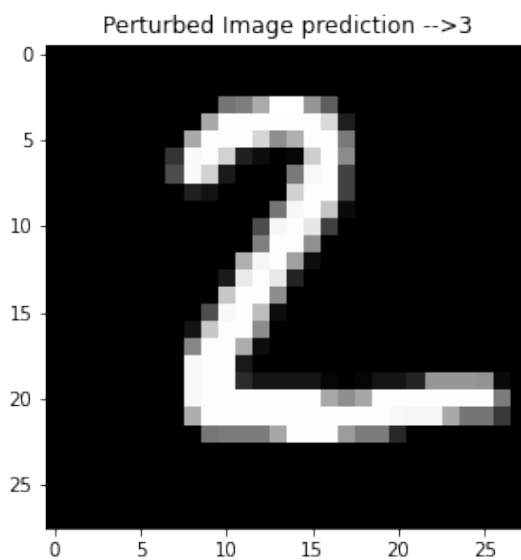

In [0]: for i in range (0,10):
        for j in range (0, len(adv_examples)):
            target, origp, adv, ex, origimg = adv_examples[j]
            if target == i:
                fig= plt.figure(figsize=(10,10))
                a = fig.add_subplot(1, 2, 1)
                a.set_title('Target: ' + str(target.item()) + ' / y_prediction w/o perturb: ' + str(y_prediction.item()))
                imgplot = plt.imshow(ex, cmap= 'gray')





                a = fig.add_subplot(1, 2, 2)
                a.set_title('Perturbed Image prediction -->' + str(adv))
                imgplot = plt.imshow(ex, cmap= 'gray')



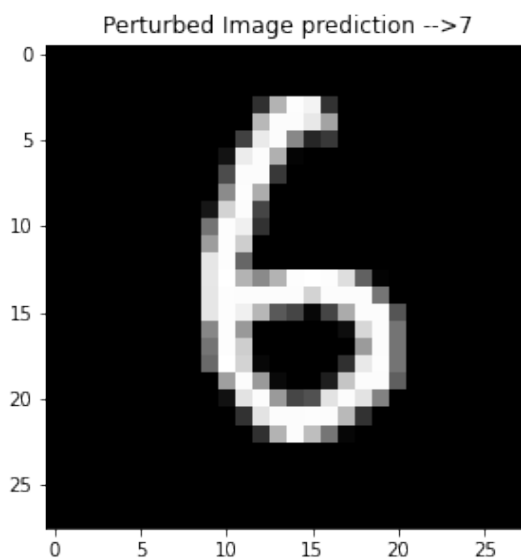

```

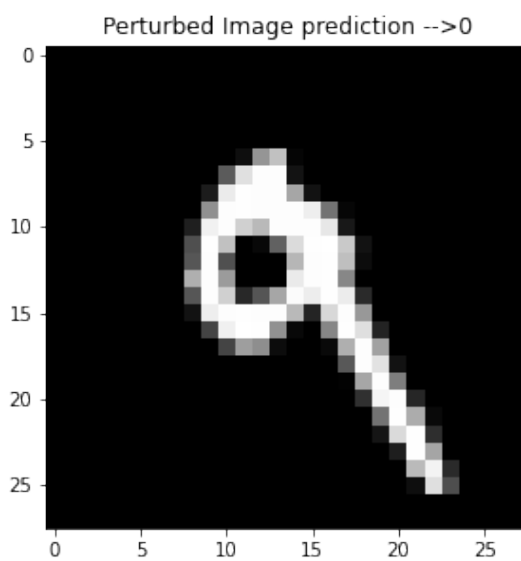


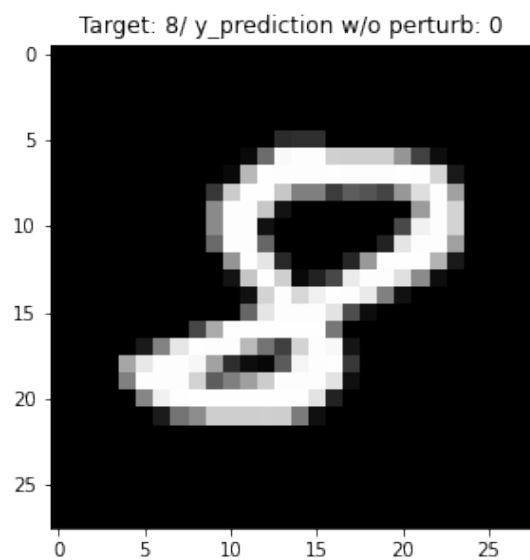
```
plt.show()  
break
```











In [0]: